

# 03 Peak calling and annotation

Grigorios Georgolopoulos

2024-07-01

## Table of contents

<b>1 Peak calling</b>	<b>1</b>
<b>2 Operations with genomic intervals</b>	<b>2</b>
<b>3 Calculating similarity between sets of intervals</b>	<b>4</b>
<b>4 Count how many peaks overlap between two sets</b>	<b>5</b>

In this chapter you will learn about what are peak calling, peak annotation, as well as identifying transcription factor binding motifs and performing motif enrichment tests.

### Learning objectives

1. Learn how to perform operations with genomic intervals
2. Compare peak sets between replicates and between samples
3. Annotate peaks
4. Map motifs to peaks and perform motif enrichment tests

## 1 Peak calling

Peak calling is a statistical method which allows us to identify regions enriched with fragments of accessible chromatin. This is a step necessary for all chromatin accessibility assays such as ATAC-seq, ChIP-seq, DNase I-seq, MNase-seq, FAIRE-seq etc.

The principle is simple. We break the genome into small windows, and we test whether the number of fragments (or reads) in a window is statistically higher than the number of fragments expected.

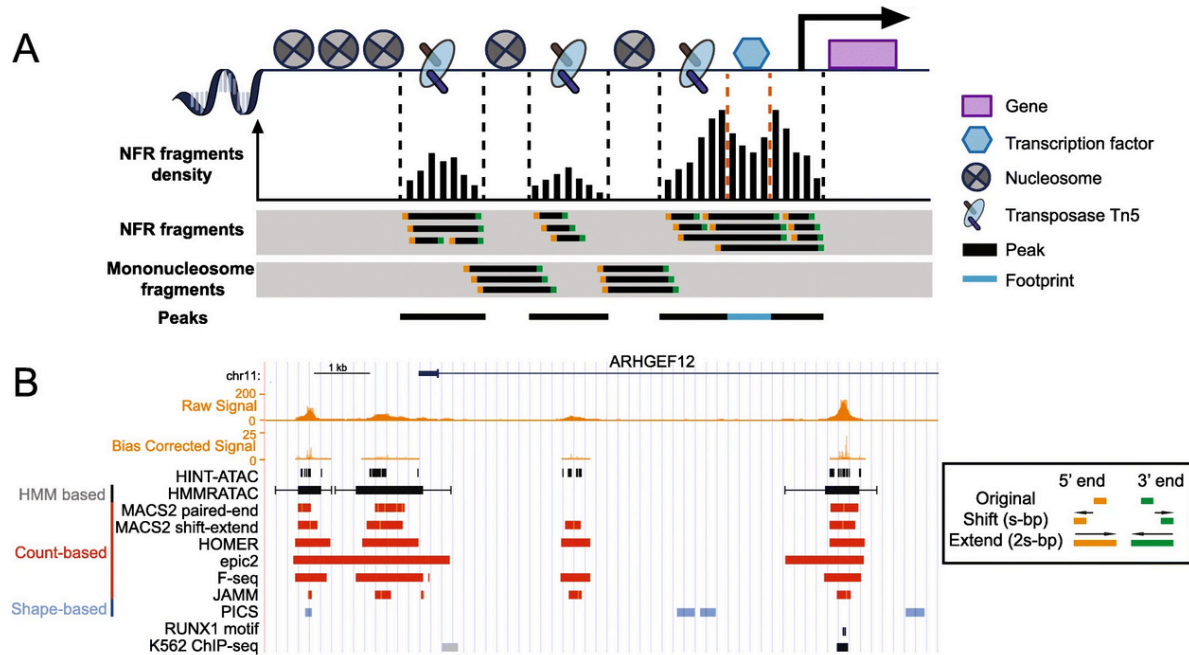


Figure 1: Peak calling overview from Yan et al., 2020

## 2 Operations with genomic intervals

One of the most common tasks in genomic data analysis is genome arithmetic. How many genomic elements in list A overlap with elements in list B? What is the average ChIP-seq signal over the regions in my list? How far is this element for the nearest TSS? How can I extract intronic regions of a gene given the exons? These are some examples that involve genome arithmetic operations. Genome arithmetic operations occur on a genomic coordinate system which requires at least two values to define an interval; a start coordinate and an end coordinate which is typically denoted like so: 1234-2345. If we operate on multiple sequences/scaffolds/chromosomes then we would need the sequence name (or chromosome name) and the notation would look like this chrZ:1234:2345.




Figure 2: An example of calculating the intersects between two sets of intervals. From bedops documentation.

There are a few established tools for performing genome arithmetic. The most popular one is [bedtools](#) which is a swiss army-knife able to perform all kinds of genomics analysis tasks with genome arithmetics being the most widely used function. Similar to bedTools there is [BEDOPS](#) which is preferred for its simpler command format and speed, although latest bedtools implementations seem to perform equally fast. For performance benchmarking information see [here](#).

Let's create a results directory first.

```
RESULTS="results/03_peaks"
mkdir -p ${RESULTS}
```

First, we need to unzip the compressed bed files

```
for GZ in data/*/*bed.gz; do
gunzip $GZ
done
```

**Before performing any interval operations we need to make sure the files are position sorted!**

BEDOPS offers a convenient function for fast sorting of genomic intervals

```
for BED in data/*/*.bed; do
echo "Sorting ${BED}..."
sort-bed $BED > ${BED}.sorted
done
```

Now let's count how many peaks are there per sample:

```
for BED in data/*/*.bed.sorted; do
wc -l ${BED}
done
```

```
149626 data/ENCSR558ZSN/ENCFF002MKC.bed.sorted
142405 data/ENCSR558ZSN/ENCFF235RAD.bed.sorted
112340 data/ENCSR977LVI/ENCFF284IBU.bed.sorted
111615 data/ENCSR977LVI/ENCFF851MGR.bed.sorted
```

According to ENCODE [ATAC-seq standards](#), for an acceptable ATAC-seq experiment we expect >100,000 peaks per sample.

### 3 Calculating similarity between sets of intervals

In this exercise we have 4 sets of peaks called from the two samples with 2 replicates each:

Sample	Info	File
ENCSR977LVI	Homo sapiens T-cell female adult (21 years), rep 1	ENCFF002MKC.bed.gz
ENCSR977LVI	Homo sapiens T-cell female adult (21 years), rep 2	ENCFF235RAD.bed.gz
ENCSR558ZSN	Homo sapiens activated T-cell female adult (21 years), rep 1	ENCFF851MGR.bed.gz
ENCSR558ZSN	Homo sapiens activated T-cell female adult (21 years), rep 2	ENCFF284IBU.bed.gz

One of the first tasks when dealing with replicates it's assessing the replicate concordance, i.e. how much the two peak sets overlap. One way to do this is by calculating the Jaccard similarity between the two sets. Jaccard similarity is calculating by dividing the intersect over the union. `bedtools` has a very convenient function to calculate this. For details on the `bedtools jaccard` function see [here](#).

```
samples=("ENCSR977LVI" "ENCSR558ZSN")
for sample in ${samples[@]}; do
files=(`ls data/${sample}/*.bed.sorted`)
fileA=${files[0]}
fileB=${files[1]}
```

```
echo "Jaccard similarity report between replicate 1 and replicate 2 for sample ${sample}:"
bedtools jaccard -a ${fileA} -b ${fileB}
echo -e "\n"
done
```

```
Jaccard similarity report between replicate 1 and replicate 2 for sample ENCSR977LVI:
intersection      union      jaccard  n_intersections
33571756          41880240          0.801613          54880
```

```
Jaccard similarity report between replicate 1 and replicate 2 for sample ENCSR558ZSN:
intersection      union      jaccard  n_intersections
44143512          57000632          0.774439          76484
```

## 4 Count how many peaks overlap between two sets

A common task is to test how many peaks are shared between two samples. This will answer questions such as how similar is the accessible chromatin landscape between two treatments/cell types/timepoints etc.

In this exercise we will evaluate how many peaks discovered from unstimulated T cells (ENCSR977LVI) are overlapping by at least 50% peaks in the stimulated T cell experiment (ENCSR558ZSN).

```
# Chose one replicate per sample
unstim="data/ENCSR977LVI/ENCFF851MGR.bed.sorted"
stim="data/ENCSR558ZSN/ENCFF002MKC.bed.sorted"

bedops -e 50% ${unstim} ${stim} > ${RESULTS}/unstim_peaks_shared_with_stim.bed

wc -l ${RESULTS}/unstim_peaks_shared_with_stim.bed
```

```
87359 results/03_peaks/unstim_peaks_shared_with_stim.bed
```

*Why didn't we try `bedtools jaccard` this time?*

*How many peaks in the stimulated set overlap by 50% the peaks in the unstimulated set?*